

Module 3

Simulating Attention in Memory
&
Learning Dot-product attention

Mini-Project 3: Attention-based retrieval from memory

Recall a word from memory based on 'cue'

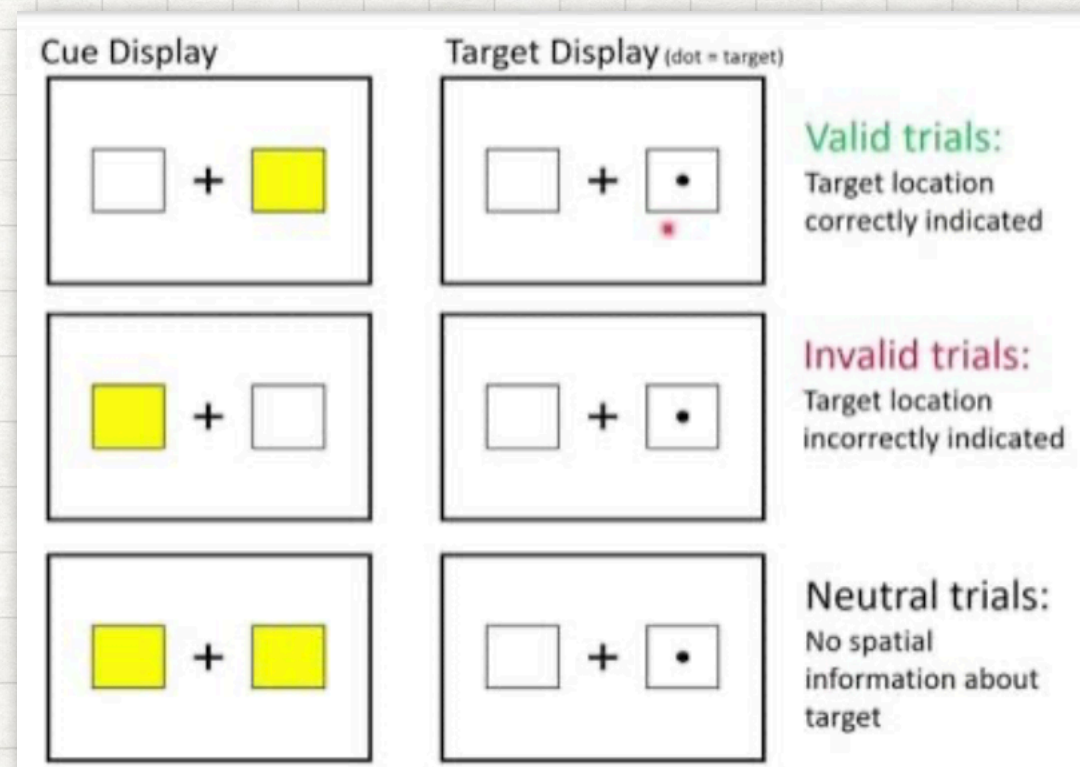
cue: valid / invalid



time


Mini-Project 3: Attention-based retrieval from memory

- * Recall depends on *match* between cue & memory (Tulving & Thomson, 1973)
- * Words are encoded in a semantic context in memory
- * Attention can be diverted to target (valid) or away from target (invalid)... Posner cueing paradigm




Mini-Project 3: Attention-based retrieval from memory

How can we simulate this memory process?




Step 1: Encode words in a high-dimensional space




Note: Similarity using cosine similarity measure

Key part of "dot-product attention"



Demo 1: Encode words in high-dimensional space



We are going to write Python code to convert:



Demo 1: Encode words in high-dimensional space


Write Python code to:

- * Create a 'WordEncoder' class. This class should:
 - * Contain an attribute called 'dim' (the dimension of vector space)
 - * Contain a method call encode()
 - * This method should take a string as input and convert it into a random vector in an n-dimensional space
 - * This random vector should lie on the unit circle

Demo 2: Find similarity between representations

Write Python code to:

- * Define a function class called 'Similarity', with functions cosine() & euclidean() that calculate the cosine & euclidean distances
- * Use one of these functions based on the initialisation of Similarity instance.



APPLE



n-dim

BREAD




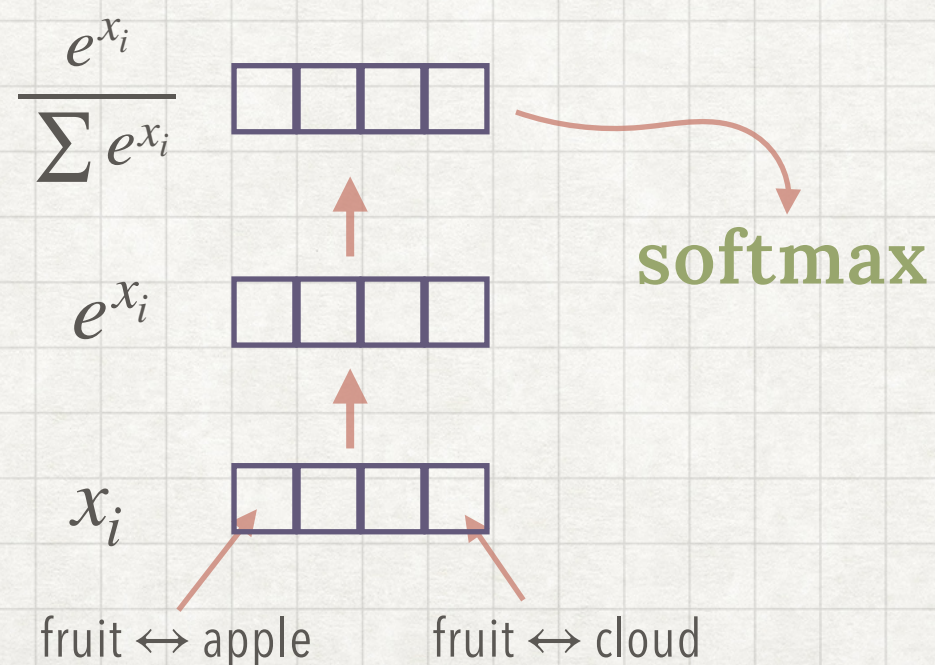
n-dim

$$\cos(\theta) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|}$$

Demo 3: Scores to probabilities

Write Python code to:

- * Create a numpy array called scores, with distances to each stored memory
- * Convert these scores to a probability distribution over memory traces



Demo 3: Pseudocode

FUNCTION softmax(scores):

1) exponentiate to amplify differences


$e_i \leftarrow \exp(z_i)$ for each z_i

2) normalize to make probabilities

total $\leftarrow \text{sum}(e_i)$

$w_i \leftarrow e_i / \text{total}$ for each i

RETURN $[w_1, \dots, w_n]$



Demo 4: Scale the attention based on temperature

Softmax function allows temperature scaling - that is how sharply focused the attention is.

When temperature is high, attention is "diffused", while when temperature is low, attention is highly focused (winner-take-all).

- * Write Python code to change the softmax function so that the probabilities are computed based on temperature.
- * Test the function for different values of temperature parameter, but same scores

$$\frac{e^{x_i}}{\sum e^{x_i}} \longrightarrow \frac{e^{\frac{x_i}{\tau}}}{\sum e^{\frac{x_i}{\tau}}}$$

Demo 3: Pseudocode

FUNCTION softmax(scores, temperature τ):

1) temperature scaling

for each s in scores:

$$z_i \leftarrow s / \tau$$

2) exponentiate to amplify differences


$$e_i \leftarrow \exp(z_i) \text{ for each } z_i$$

3) normalize to make probabilities

$$\text{total} \leftarrow \text{sum}(e_i)$$

$$w_i \leftarrow e_i / \text{total} \text{ for each } i$$

RETURN $[w_1, \dots, w_n]$



Demo 5: Compute Context vector via Dot-product Attention

Let us put it all together and parallelise the attention mechanism.

Write a class called DotProductAttention with a function attend() that

1. computes the scores for all memories in parallel
2. Computes the attention weights for each memory
3. Computes a "context" based on aggregated memory recalled
4. Returns this context and attention weights

Test this function for memories: ['apple', 'bread', 'cloud', 'drum', 'eagle'] & query 'apple'

Cue → Query

Memory index → Key

Memory rep → Value

Demo 5: Compute Context vector via Dot-product Attention

```
CLASS DotProductAttention(temperature  $\tau$ ):
```

```
    METHOD attend(query q, keys K, values V):
```

```
         $d \leftarrow$  dimension of q
```

```
        # 1) compute similarity scores (dot products)
```

```
        for each key  $k_i$  in K:
```

```
             $score_i \leftarrow q \cdot k_i$ 
```

```
        # 2) convert scores to attention weights via softmax
```

```
         $w \leftarrow \text{softmax}([score_1, \dots, score_n], \tau)$ 
```

```
        # 3) compute context vector as weighted sum of values
```

```
         $c \leftarrow \sum_i (w_i * v_i)$ 
```

```
    RETURN w, c
```

Cue \rightarrow Query

Memory index \rightarrow Key

Memory rep \rightarrow Value